

Comparison-Page Chain — Gap Analysis

Fixes (what's broken) · Enhancements (what could be better) · Efficiency (same quality, less cost) · Additions (what's missing). Every finding pinned to a repo artifact, task record, or live-verified fact.

Version: v1 · **Date:** 2026-06-11 · **Companion to:** Production Map v1 (phase letters 0, A–K refer to it) · **Verified today:** registry · TASK-245 close record · OFF remediation brief · prod-sync brief · drafted CI · factory artifact coverage

Context that frames everything: the biggest structural defect — production serving from a repo **81 commits behind** where all work happened — was found and **fixed today** (TASK-245: bari.digital now deploys from Barint master; 9/9 routes live-verified, 0 OFF refs). This analysis covers what that fix exposed and what remains.

1 · Fixes — What Is Broken Right Now

P0 consumers see it today **P1** wrong but not yet visible **P2** hygiene

FIX-1 · Live yogurts = the thin interim shelf; the replacement is ready, parked behind 3 gates

P0 Phase J **OWNER = CRITICAL PATH**

- **Consumers today:** v3 — 11 products, **0 of 11 with ingredient data**, all confidence "partial/under review," high-protein shelf bias (plain/bio/Greek/flavored shelf absent). Plus **"NOVA" 9x in live consumer prose** (banned jargon, logged at TASK-245 close).
- **The fix exists:** run_yogurt_006 / v4 — 88 products, 0 errors, all red-team findings resolved (TASK-249/250 RETURNED).
- **Held by exactly 3 gates:** ① **your Ruling-3 sign-off** (two published grades change: 35/D→35/E, 50/C→50/D — tripwire) · ② Content's ceiling-compression caveat copy · ③ QA baseline freeze on run_006.

Action: highest-leverage consumer-facing move available — work done, gates only. Your ruling is the critical path.

FIX-2 · QA passed a corpus red-team demolished — the QA gate tests traceability, not truth

P1 Phases B–E **QA** **DATA**

- run_005 cleared QA; red-team then found ingredient text corrupted on **67 of 89 products** + a protein=190 value no gate blocked. The QA protocol checks coverage/traceability/versions — nothing checks whether parsed values are *plausible* or parsed text is *ingredients at all*.
- Truth-checking happens only at Phase F — the most expensive place (after scoring, copy, packaging → full regen; yogurt took 6 runs).
- The fix is half-built and stuck: **TASK-218** (scrape QA standard) sits RETURNED; the run_006 machinery (`macros_plausible`, `ingredient_text_quality`, disclaimer-strip) protects **one category only**.

Action: close TASK-218 → blocking Stage-B/C gate; lift run_006 checks into shared BSIP1 core.

FIX-3 · Five RETURNED tasks queued at the close gate — verification is the single-threaded bottleneck P1 Phase K ORCHESTRATOR / CC

- **Queue:** TASK-217 (juice NOVA-1 gate) · 218 (scrape QA standard) · 241 (salty 12-rescue) · 249 + 250 (yogurt v4). Finished work, invisible to the roadmap until closed; whatever they `block` stays shut.

Action: batch close-readiness session now; route mechanical verification (`file-exists` / `counts` / `grep-zero`) to the cheap-LLM lane (CC v4.1) — CC spends judgment only on judgment.

FIX-4 · Known engine defects — all registered, none landed P1 Phase D NUTRITION

Defect	Task	State	Risk
Fermentation keyword bonus uncapped (TRIM Path A)	TASK-246	IN_PROGRESS	Keyword stuffing can ride toward A/S
Absence-as-zero — missing nutrient scores as 0, cross-category	TASK-236	BLOCKED	Punishes honest missing data — contradicts "unknown is acceptable"
Juice NOVA-1 floor fires for reconstituted-from-concentrate	TASK-217	RETURNED	Unearned whole-food floor
Router v2: no bread/cracker archetype; seed-topped bread → WFF	—	unowned	Misroute → wrong calorie table → wrong score

Action: prioritize TASK-236 — a *philosophical* defect (the engine violates its own data-honesty principle) touching every category's confidence story.

FIX-5 · Registry / reality drift P2 Phase K CC

- TASK-237 (salty OFF re-source) IN_PROGRESS on paper — but salty v4 with 29 real products **shipped to production today**.
- Memory index says "TASK-189 open (granola sodium)" — it **closed 2026-06-05** with `BARI_SODIUM_CEREAL` implemented.
- Four blog tasks (203–206) parked IN_PROGRESS indefinitely — dispatch or park explicitly.

Action: one CC drift sweep ≈ one hour.

FIX-6 · "NOVA" jargon ruling pending on two more live pages P2 Phase G CONTENT

Maadanim (86 hits) and hummus (70 hits) show "NOVA" in metric cells — designed column or leak? Routed at TASK-245 close; not yet ruled. Designed = keep + register exception; leak = template fix.

2 · Enhancements — Works, but Could Be Better

ENH-1 · Generalize the yogurt honesty gates to every category

Phases B-C

DATA

Disclaimer-strip + `macros_plausible` + `ingredient_text_quality` is the best data-quality work in the pipeline — protecting exactly one category. Lift into `03_operations/bsip1/core/`; the shared `bsip0_gate.py` + 30 tests already exist (untracked) — finish, commit, wire.

ENH-2 · Finish the verdict-copy rollout — 8 legacy categories pending

Phase G

CONTENT

The ruled standard (verdicts name calorie density + the real fired driver; sodium = displayed fact only) is live in new categories only. One Content pass per legacy page; zero data changes.

ENH-3 · Image story after de-OFF

Phases B-J

QA

DATA

The OFF purge correctly nulled contaminated images → live placeholders in several categories. Backfills ran (TASK-208/243 closed). Needed: one sweep — % real scrape-sourced images per live category + a stated SLA ("no launch under X%; placeholders acceptable, OFF never").

ENH-4 · Wire the two finished-but-idle quality frameworks into the chain

Phases G-D

ORCH

CONTENT

- **Hebrew Golden Eval** (TASK-220, built) → automatic Phase-G gate on every copy batch; mechanizes read-every-string.
- **Observability event schema** (TASK-219, built) → pipeline runs emit events; "what ran, when, which config hash" stops living only in per-run report files.

ENH-5 · Red-team earlier *and* at the end

Phases C-F

RED-TEAM

Keep Phase F as the full adversarial audit; add a **30-minute "data sanity" pass right after BSIP1** (sample 10 products: raw text vs parsed output, macro plausibility). Run_005's disaster would have cost one agent-hour instead of a full regeneration cycle.

3 · Efficiency — Same Quality, Less Time & Cost

EFF-1 · Shift truth checks left — the 6-run lesson

Yogurt's cost structure: errors injected at B/C, discovered at F → B-to-F reruns. Every plausibility assertion moved from F to C converts a multi-session regen into a minutes-level parse fix. **The same checks, earlier, are ~10× cheaper.**

EFF-2 · Commit the drafted CI — it's written and has never run

DATA

FRONTEND

`.github/workflows/barint_ci.yml` already defines the right gates (next build · lint · corpus validation · BSIP0 nutrition tests · 64 enricher tests · router + golden-corpus regressions · BSIP0 gate tests) — but the folder is **untracked; no PR has ever been CI-checked**. One commit makes the frozen-invariant and build gates automatic.

Add one missing job: the OFF-string sweep (`openfoodfacts|off_api` over `bari-web/src/data/`) — today that Phase-J check is manual.

EFF-3 · Finish the shared packaging core — TASK-233F, 10 generators pending

DATA

Each category has a bespoke builder (`build_yogurts_frontend_v4.py` , `_v006.py` ...) — every honesty gate and format change re-implemented per category. That is exactly how single-category safeguards happen. One core = write once, all categories inherit.

EFF-4 · Reconcile the factory skill with practice

The skill mandates per-stage artifacts; only 4 categories have them (cereals, cheese_spreads, juices, hard_cheeses) — yogurt, the newest and most scrutinized run, produced none. Either the shared runner **emits them automatically** (preferred; free once EFF-3 lands) or the skill trims to artifacts that earn their keep. A spec the flagship run ignores trains agents to ignore the spec.

EFF-5 · Unblock the close gate — cheap-LLM verification lane

CC

Mechanical claim-verification (file exists, counts match, grep clean, build green) is zero-judgment work → Stupid-LLM lane with idiot-proofed prompts; CC reviews the evidence table. Target: RETURNED→CLOSED in hours, not days.

EFF-6 · Parallelize copy (G) and packaging (H) after score freeze

They share one dependency: frozen scores + traces. After QA freeze, dispatch Content and Data in parallel; H consumes G's strings at the end. ≈ one session saved per category.

EFF-7 · Name the Shufersal-VPN dependency as an operational gate

Yogurt re-acquisition waited on Israeli-IP VPN availability — an undocumented bottleneck for every future scrape. Cheap fix: a scrape-window runbook (when VPN is up, which scrapes queue, who flips it) so Phase B never silently stalls.

4 · Additions — What Doesn't Exist Yet

ADD-1 · Post-deploy production smoke test (automated) — must-have

The prod/local split survived for months because **nothing watched the live site**. TASK-245 fixed the topology; CC verified the flip manually — but no recurring check exists. Minimum: a scheduled job hitting every `/hashvaot/*` route asserting HTTP 200, 0 OFF refs, expected data-version marker. **The regression alarm for the whole "what's actually live?" failure class.**

ADD-2 · Live-state manifest: route → data version → run ID → baseline — must-have

Gap 2 of the prod-sync brief was exactly this absence — enumerating the re-point payload required hand-diffing repos. A small machine-readable manifest, written by packaging and asserted by the smoke test, makes "what is live?" a lookup instead of an investigation.

ADD-3 · BSIP0 scrape QA standard as a blocking gate — must-have

TASK-218's deliverable, closed and wired into Phase B for every future scrape (institutional form of FIX-2). The weakest link in the chain is between retailer HTML and BSIP1 — it deserves a standing gate, not per-category heroics.

ADD-4 · Data-freshness policy per category — must-have

No policy says when a category must re-verify (reformulations, recalls, price moves). Even coarse — "re-scrape + re-verify every N months; show the date; stale categories get a banner" — converts an unbounded liability into a scheduled cost.

ADD-5 → ADD-9 · Nice-to-have

Addition	Value
ADD-5 Consumer analytics loop	The 15–20s comprehension metric is unmeasured in prod; traffic/scroll/filter data lets Product sequence the roadmap on demand, not intuition.
ADD-6 Editorial regression suite	Banned-phrase + jargon linter (9 banned phrases, "NOVA", score-presentation rules) at build time — FIX-6 becomes impossible to ship.
ADD-7 Price display	<code>il_prices</code> client exists (TASK-170). Price-per-100g next to scores = major consumer value; needs its own freshness rule. Product-gated.
ADD-8 Category health dashboard	Per category: corpus age, last run, gates passed, live version, open findings. Command center tracks <i>tasks</i> ; nothing tracks <i>categories</i> as first-class objects.
ADD-9 Cross-market disclosure	Wave 3 backlog, already scoped (D5 annotate-only, never a score move). Parked correctly; listed so it isn't lost.

5 · If You Only Do Five Things

#	Action	Type	Why first
1	Rule on yogurt Ruling 3 → ship v4 (88 products) over the thin 11-product live shelf	FIX-1	Biggest consumer win; work done; you are the critical path
2	Batch-close the RETURNED queue (217/218/241/249/250) via the cheap-verification lane	FIX-3 + EFF-5	Unlocks everything queued behind finished work
3	Commit the CI workflows + add the OFF-sweep job	EFF-2	One commit; 5 manual gates become automatic forever
4	Post-deploy smoke test + live-state manifest	ADD-1 + ADD-2	Closes the failure class behind the worst defect in project history
5	Generalize yogurt honesty gates into shared BSIP1 + adopt TASK-218	ENH-1 + ADD-3	Makes the 6-run lesson permanent for every future category