BARI · ברי

PARTNERSHIP BRIEFING — CONFIDENTIAL

The trustworthy shelf for food & supplements.

An evidence-graded comparison platform for the Israeli consumer — built on a deterministic scoring engine, a disciplined data pipeline, and a multi-agent production system.

Deterministic scoring engine | Direct-scrape data only | 13 live categories | Multi-agent factory

Prepared for: Collaboration exploration — nutritionist & engineering partner
Prepared by: Tom Bar Haim, Founder & Owner, Bari
Date: 21 June 2026 · **Contents:** Architecture · Engine · Evidence · Agents · Business plan

BRIEFING CONTENTS

What's in this memo

Two readers in mind: a clear strategic picture for the founder, and enough technical depth to answer an engineer's questions on the spot. Technical deep-dives are flagged FOR THE ENGINEER throughout.

01 — EXECUTIVE SUMMARY

What Bari is, in one breath

Bari turns the noise of a supermarket shelf into a single, defensible verdict. For any food or supplement category, it answers the only question a shopper actually has: *“of the things in front of me, which is genuinely better, and why?”*

13 LIVE CATEGORIES	**10** SCORING DIMENSIONS	**0** INVENTED DATA POINTS	**6** AI LANES IN THE FACTORY

The gap we close

Existing tools fail the consumer in one of two ways. Government databases and apps like Open Food Facts are **crowd-sourced and unreliable** — wrong barcodes, missing fields, contributed errors. Brand marketing is, by definition, **not neutral**. Nobody gives an Israeli shopper a *trustworthy, explained, category-relative ranking* in Hebrew. That is the whole of what Bari does.

WHAT MAKES IT CREDIBLE

- **Deterministic engine.** Scores come from auditable Python, not an AI guessing — every score reproduces from a trace file.
- **Direct-scrape data only.** If we didn't read it off the real retailer/label, the field is empty and the page says so. We never fill gaps with crowd data.
- **Evidence registry.** Every scoring rule cites a registered evidence code (EV-###); rules can't accumulate without justification.
- **Two-gate content.** No consumer sentence ships without a content pass *and* an adversarial red-team pass.

WHY IT SCALES

- **A factory, not a craft.** Each category is the output of one repeatable pipeline — not a hand-built page.
- **One engine, one path.** Change a scoring rule once and every category re-flows through the same machinery.
- **Multi-agent labour.** Scraping, enrichment, copywriting, QA and research run across specialised AI lanes coordinated by an orchestrator.
- **Hebrew-native.** The whole stack reasons over Hebrew ingredient text — a real moat in this market.

> **THE ONE-LINE POSITIONING**
>
> “Bari is the **category-relative, evidence-graded shelf** for Israeli food and supplements — the only place a shopper gets an honest verdict, with the reason shown, in their own language.”

What's live today

Bari is a working Hebrew web platform (bari.digital) organised around **comparison pages**. Each page takes a real shelf of products, scores them, and presents a ranked, explained verdict shoppers can read in 15–20 seconds on a phone.

Categories shipped through the pipeline

CATEGORY	STATUS	CATEGORY	STATUS
Milk & alternatives	LIVE	Hummus & spreads	LIVE
Bread	LIVE	Juices	LIVE
Brined cheeses	LIVE	Hard cheeses	LIVE
Yellow / sliced cheese	LIVE	Breakfast cereals	LIVE
Granola	LIVE	Cookies & coffee biscuits	LIVE
Cakes & hard cookies	LIVE	Salty snacks	LIVE
Protein & snack bars	LIVE	Supplements (magnesium pilot)	IN BUILD

All live categories pass a structural conformance gate (12/12 verified 2026-06-18) — meaning a change to the scoring engine re-flows cleanly through every one of them.

What a comparison page gives the shopper

THE VERDICT

A numeric score (0–100) and grade (S–E) per product, ranked within the category. No color-coded spin — the number and the grade carry it.

THE REASON

A two-line human verdict per product: where it stands, why, and the catch. The engine's real driver is named (e.g. calorie density, sodium, fat source).

THE HONESTY

A category caveat box, confidence states when data is thin, and an explicit "data could not be retrieved" rather than a guessed value.

DESIGN DISCIPLINE

Every page conforms to a single golden template — frozen pixel geometry, a fixed set of canonical components, Hebrew RTL correctness, and WCAG contrast checked by a vision-grounded design critic. We are deliberately in a *conformance* phase, not a redesign phase: consistency is the feature.

The data pipeline

A product becomes a scored, published page by moving through three stages we call **BSIP0 → BSIP1 → BSIP2** (Bari Structured Intelligence Pipeline). Raw retail data in; a defensible verdict out.

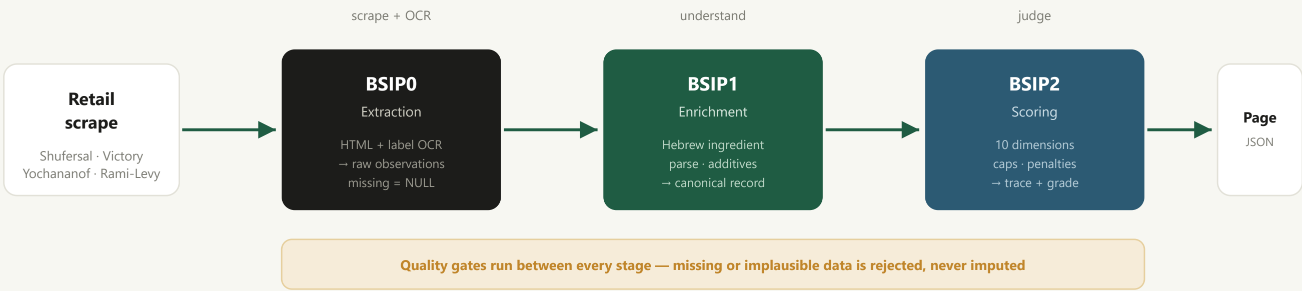

Figure 1 — The BSIP pipeline. The engine never reads external databases; it only sees what the scrape captured.

BSIP0 · EXTRACTION

A 4-stage scraper (discover → approve → scrape → audit) pulls product HTML from up to four Israeli retailers; an Azure-based OCR pipeline reads physical label images. Output: raw per-product observations. **If a field isn't found, it stays NULL** — no imputation, ever.

BSIP1 · ENRICHMENT

Python parses the Hebrew ingredient list: detects white vs. whole-grain flour, fermentation markers, additive classes and burden count, protein source, and food-structure signals. Assigns a trust level (high/medium/low) and writes a single canonical JSON record per product.

BSIP2 · SCORING

The deterministic engine extracts 50+ features, scores 10 dimensions, applies caps, penalties and family budgets, and writes a full *trace* file. The website consumes only this generated JSON — the score on screen always reproduces from the trace.

FOR THE ENGINEER

Stages are decoupled file-to-file contracts: each writes JSON the next stage loads. No stage calls a model to *invent* a value. AI is used for *parsing assistance and copy*, never to produce a nutrition number or a score. That separation is what makes the output auditable and reproducible.

04 — THE SCORING ENGINE

How a score is computed

This is the heart of Bari and the part most likely to draw technical questions. The engine (internally "BSIP2 v0.4.1") is **pure deterministic Python** — same input, same score, every time. No randomness, no LLM in the scoring path.

THE 10 DIMENSIONS (EACH 0–100)

DIMENSION	WT	READS
Processing quality	15%	NOVA level, additive burden
Nutrient density	15%	Protein, fiber, micros / 100g
Calorie density	15%	Calories vs. category norm
Glycemic quality	12%	Sugar load, carb make-up
Protein quality	10%	Quantity + source
Additive quality	10%	Stabilizers, emulsifiers, sweeteners
Fat quality	8%	Sat-fat proportion, fat source
Satiety support	6%	Protein + fiber + structure
Regulatory quality	5%	Israeli red-label warnings
Whole-food integrity	4%	Proximity to whole food

Weights are prototype values that sum to 1.0; public calibration is a separate, deliberate phase.

THE 6-STAGE RESOLUTION

1 · Feature extraction
50+ features; missing fields recorded, never filled

2 · Dimension scoring
10 dimensions → weighted base score

3 · Guardrails
Vetoes, hard caps, soft penalties, floors

4 · Hyper-palatability
Fat-sugar / fat-salt combos, with budget cap

5 · Concern coordination
Stops one concern penalising twice

6 · Final resolution
Caps + penalties + confidence ceiling → grade

trace.json · score 0–100 · grade S–E

> **FOR THE ENGINEER — KEY DESIGN CHOICES**
>
> **Caps vs. penalties:** caps are binding ceilings (NOVA-4 → cap, multiple red labels → cap); the most restrictive cap wins. Penalties are subtractive and bounded by per-family "budgets" so a single root concern (sugar, sodium, fat) can't stack penalties across rules. **Confidence ceilings** bound the score when data is thin. Everything is in one `constants.py`; a `router_v2` assigns category before scoring; outputs are emitted by a `trace_writer`. Grades: S 90+, A 80, B 65, C 50, D 35, E 0.

05 — EVIDENCE & INTEGRITY

Why the numbers can be defended

A score is only as trustworthy as the discipline behind it. Bari runs four hard rules that exist precisely so a verdict can be defended publicly — to a brand, a regulator, or a journalist.

1 · THE OFF BAN (PROJECT-WIDE, PERMANENT)

Open Food Facts and any crowd-sourced source is **banned for every field, forever** — nutrition, ingredients, names, barcodes, images. The only source is the direct product scrape. If it wasn't scraped, the field is NULL and the page says “data could not be retrieved.” *“Unknown is acceptable; OFF is not.”*

2 · THE EVIDENCE REGISTRY (EV-###)

Every scoring rule and adjustment cites a registered evidence code. New rules can't accumulate silently — each needs a justification entry, an activation scope, and a rollback plan. This stops “rule sprawl” and keeps the engine governable.

3 · CITATIONS DISCIPLINE

Every research, nutrition or consumer claim names its source inline — a URL, a document, a trace, or “direct product scrape.” Vague provenance (“an official food source”) is banned at the QA gate. No claim ships without a traceable origin.

4 · NO INVENTED DATA

The engine never fabricates a product, nutrition value or ingredient. Missing-data products are discarded rather than guessed or “temporarily” filled. Honesty over coverage — a smaller true shelf beats a larger fake one.

How a scoring change is allowed to happen

New evidence (EV-### registered) → **Behind a flag** (default OFF, scoped) → **Re-flow all** (conformance gate) → **Red-team** (defend every move) → **Owner merge** (never auto-deploy)

Figure 2 — No score changes published without evidence, a re-flow across all categories, an adversarial pass, and a human merge.

WHY THIS MATTERS FOR A PARTNERSHIP

A nutritionist's name on a platform is only as safe as that platform's worst claim. Bari is built so that *every* claim is traceable and every score is reproducible — the integrity discipline is the asset, not the friction.

What it's built on

Two clean concerns in one monorepo: a **Python data/scoring workspace** at the root, and a **Next.js website** as a self-contained subtree. They meet at exactly one interface — generated JSON.

BACKEND · DATA & SCORING

LAYER	TECH
Language	Python 3 (shared venv)
Scraping	Custom 4-stage retailer scrapers
OCR	Azure-based label extractor
Enrichment	Hebrew ingredient parser + additive classifier
Scoring	Deterministic engine (13 source modules)
Output	Per-product `trace.json` + frontend JSON
Tests	pytest suites (e.g. 64 enricher checks)

FRONTEND · THE WEBSITE

LAYER	TECH
Framework	Next.js (React, App Router)
Language	TypeScript
Styling	Tailwind CSS, design tokens
Layout	Hebrew RTL, mobile-first
Data in	Generated JSON in `src/data/comparisons/`
View model	Strict VM contract: UI never sorts / rounds / interprets
Hosting	Vercel-class static/SSR deploy

Repository shape

FOLDER	ROLE
01_framework	The “intelligence framework” — scoring theory, concept definitions, governance docs. No code.
02_products	Category-first workspaces — raw sources, observations, canonical records, scored outputs per category.
03_operations	The factory — scrapers, OCR, the scoring engine, batch runners, page generator.
bari-web	The Next.js website (git subtree). Consumes generated JSON only.
tasks / .claude	The Agent OS — task registry, agent definitions, routing law, skills.

FOR THE ENGINEER — THE SEAM

The website is intentionally “dumb”: it renders a view model and never re-computes anything. All intelligence lives upstream in Python and is frozen into JSON artifacts before the site sees it. This means the frontend and the engine can be worked on independently, and the published numbers can't drift from the trace that produced them. Categories re-publish by re-running the generator, not by editing pages.

How the work actually gets done

Bari is produced by a **multi-agent system** — a coordinated set of AI lanes, each with a defined job, run by an orchestrator off a single task registry. Think of it as a small, specialised studio that never sleeps.

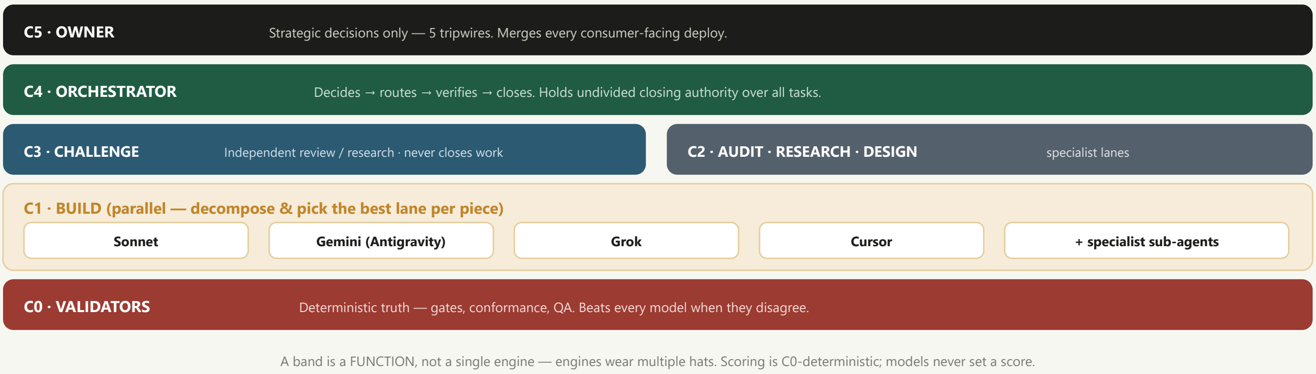

Figure 3 — The C0–C5 band model. Humans decide strategy and merge; models build and review; validators are the final arbiter of truth.

SPECIALIST AGENTS

Named roles — Data, Nutrition, Research, Product, Frontend, Design, Marketing, and an Adversarial QA / Red-Team gate — each with its own remit and tools.

ONE REGISTRY, ONE TRUTH

Every tracked job is a `TASK-###` file with a strict lifecycle. The orchestrator verifies each claim against artifacts before closing — no “trust me, it's done.”

TWO-GATE CONTENT

No consumer sentence ships without a Content pass *and* an adversarial Red-Team pass. Both sign off, every time, every category.

Done · Doing · Going

Bari is past “does it work” and into “make it broad and make it pay.” Here is the honest state.

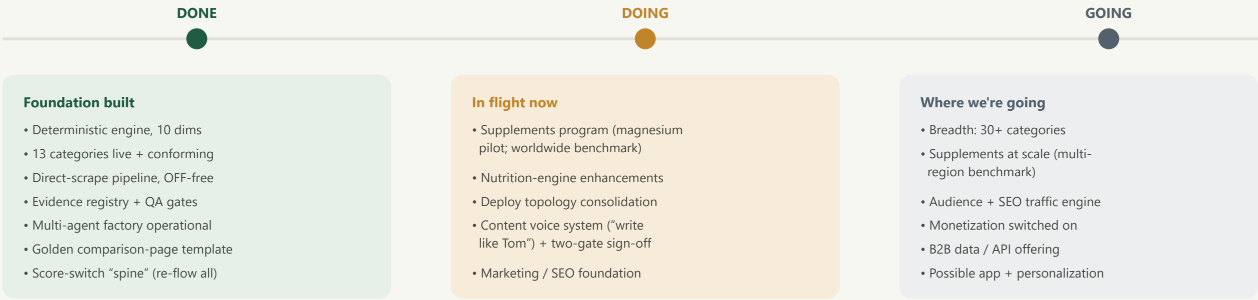

Figure 4 — Strategic timeline. The engine and factory are built; the focus shifts to breadth, audience and revenue.

13→30+	**1**	**$0**
CATEGORY TARGET	ENGINE, ALL CATEGORIES	REVENUE SWITCHED ON YET

THE HONEST READ

The hard, defensible part — a trustworthy engine and a repeatable factory — is **done**. What remains is largely *volume and distribution*: more categories, more traffic, and turning on revenue. That is exactly the phase where the right partners add leverage.

09 — BUSINESS PLAN

Market & monetization

Bari starts as a high-trust Israeli consumer destination and compounds into a data & intelligence business. The audience is the moat; the engine is the product nobody else can cheaply copy.

The market

WHO

Israeli grocery & supplement shoppers — health-curious, mobile-first, Hebrew-speaking. Secondary: dietitians, trainers and clinics who need a neutral reference.

WHY NOW

Rising health awareness, distrust of brand claims, weak/unreliable local data, and AI that finally makes per-category, explained scoring economical to produce.

EDGE

Hebrew-native reasoning, an auditable engine, and an integrity reputation. Hard to fake, expensive to catch up to, defensible in public.

Four revenue streams (sequenced, not simultaneous)

STREAM	WHAT	WHEN	WHY IT FITS BARI
1 · Audience / affiliate & retail referrals	Monetize buying intent at the moment of comparison — retailer / e-grocery referral, sponsored-but-labelled placements that never touch the score.	Near-term	Natural: shoppers arrive deciding what to buy. Firewalled from scoring.
2 · Premium consumer	Personalization, saved shelves, alerts, deeper category guides, “build my basket” tools.	Mid	Recurring revenue once audience & habit exist.
3 · B2B data & API	License the scored corpus / scoring API to brands (benchmarking), retailers (private-label QA), apps & clinics.	Mid–late	The engine + clean corpus is a sellable asset; high margin.
4 · Professional / white-label	A nutritionist- / clinic-facing tier: branded comparison tools, patient-facing references, supplement vetting.	Partnership-led	**Directly relevant to this meeting.**

Revenue mix as it matures

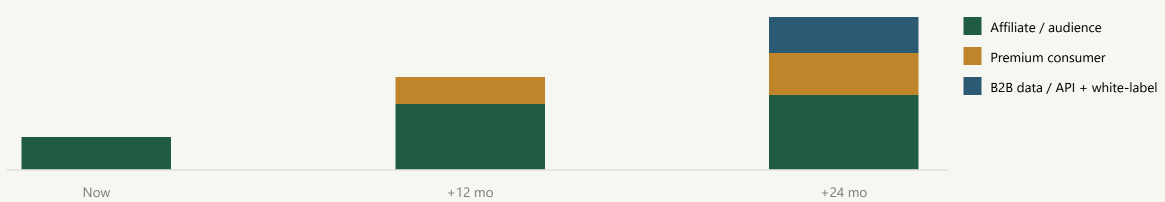

Figure 5 — Illustrative revenue mix. Audience first; premium and B2B layer on as traffic and corpus compound. (Directional, not a forecast.)

GUARDRAIL — NON-NEGOTIABLE

Money never touches the score. Sponsorship, affiliate and B2B revenue sit *outside* the scoring firewall — the day a paid relationship can move a grade, the trust asset is gone. This rule is what makes every revenue stream above safe to pursue.

How we expand

Growth has two axes: more **shelf** (categories) and more **people** (audience). The factory makes the first cheap; SEO and trust make the second compounding.

AXIS 1 — CATEGORY BREADTH

Each new category is one pass through the same pipeline. Marginal cost falls every time we add one because the engine, components and gates are shared. Sequencing favours: high search demand, clear quality spread between products, and clean scrapeable data.

- Foods: oils, pasta & grains, frozen, sauces, baby food, plant-based
- Supplements: magnesium → vitamin D, omega-3, protein, multivitamins
- Each category = a cluster of long-tail Hebrew search queries

AXIS 2 — AUDIENCE & DISTRIBUTION

Comparison pages are durable SEO assets — they answer "which X is better" queries that recur forever. Layer content marketing, a newsletter, and professional referrals (dietitians citing Bari) on top.

- SEO: own the Hebrew "which is healthiest" query space
- Authority: integrity reputation → press & professional citation
- Retention: newsletter, alerts, personalization later

The expansion flywheel

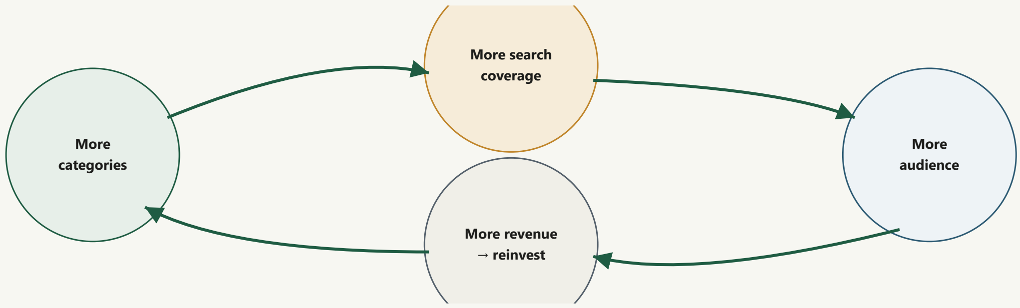

Figure 6 — Each category adds search surface → audience → revenue → funds more categories. The engine makes the loop cheap.

COSTS ARE MOSTLY FIXED

The engine and factory are built. New categories are incremental compute + review, not new infrastructure.

DEFENSIBILITY GROWS

The scored corpus, the Hebrew tooling and the integrity reputation all compound — late entrants can't shortcut them.

OPTIONALITY

Same asset supports consumer, B2B, and professional/white-label without rebuilding the core.

Where a partnership could fit

They are building something adjacent. The interesting question isn't "who wins" — it's "what does each side bring that the other can't cheaply build." A few honest framings to explore in the room.

WHAT BARI BRINGS

- A working, auditable scoring engine & data pipeline
- 13 live categories + a repeatable factory to add more
- Hebrew-native ingredient reasoning & integrity discipline
- A multi-agent production system that scales output

WHAT A NUTRITIONIST + ENGINEER COULD BRING

- Clinical credibility & a professional/patient channel
- Domain validation of scoring philosophy
- Engineering capacity to accelerate breadth
- Distribution into clinics, trainers, dietitian networks

Models worth discussing (keep options open)

MODEL	SHAPE	WATCH-OUT
Content / advisory	Nutritionist validates & co-signs methodology; named scientific authority.	Light, low-risk, fast to start.
White-label / professional tier	Bari engine powers their clinical-facing product; revenue share.	Keep the scoring firewall & brand boundaries clear.
Merge / joint venture	Combine teams and roadmaps.	Diligence: IP, equity, control of the scoring philosophy.

> **PROTECT BEFORE YOU SHARE**
>
> The engine, the corpus and the Agent OS are the crown jewels. In a first meeting, share the *vision and the discipline* generously; share engine internals, thresholds and the corpus only under a mutual NDA. The integrity story sells the partnership — the source code doesn't need to be on the table yet.

Likely engineer questions — quick answers

"Is the score an AI guess?"
No. Scoring is deterministic Python with a full trace per product; the same input always yields the same score. AI assists with parsing and copy, never with setting a number.

"Where does the data come from — is it reliable?"
Direct scrapes of up to four Israeli retailers, plus label OCR. Crowd sources (Open Food Facts) are banned entirely. Missing fields stay empty; we never impute.

"How do you stop the engine becoming an unmaintainable pile of rules?"
Every rule cites an evidence code, ships behind a flag with a rollback, and must re-flow all categories through a conformance gate before publish.

“What's the stack and how do you deploy?”

Python pipeline → generated JSON → Next.js/TypeScript/Tailwind site (Hebrew RTL). The site only renders a view model; it never re-computes. Categories republish by re-running the generator.

IN ONE SENTENCE

The engine and the factory are built and honest. The opportunity now is breadth, audience and revenue — and that is where the right partner compounds value.

Tom Bar Haim — Founder & Owner, Bari

tbarhaim@gmail.com

Prepared 21 June 2026 · Confidential — for partnership discussion only