

STRATEGIC BRIEFING — CONFIDENTIAL

# The trustworthy shelf for food & supplements.

An evidence-graded comparison platform for the Israeli consumer — a deterministic scoring engine, a disciplined data pipeline, and the go-to-market to turn it into a trusted national reference.

499 products scored

13 live categories

Deterministic engine

Only verified data

**Prepared by:** Tom Bar Haim — Founder & Owner, Bari

**Date:** 21 June 2026 · **Inside:** Product · Engine · Data integrity · The team · Market · Go-to-market

CONTENTS

# What's in this briefing

A complete picture of Bari today and where it is going — the product, how it actually works, why the numbers can be trusted, the team that produces it, the market, and a phased plan to become a national reference.

|    |                                                                        |       |
|----|------------------------------------------------------------------------|-------|
| 01 | Executive summary — what Bari is & why it's different                  | 3     |
| 02 | The product — what's live, and what it tells a shopper                 | 4     |
| 03 | How a product becomes a score — the pipeline                           | 5     |
| 04 | The engine — how a score is computed (and why it doesn't double-count) | 6     |
| 05 | Why the numbers can be trusted                                         | 7     |
| 06 | How it's built                                                         | 8     |
| 07 | The team that produces it                                              | 9     |
| 08 | The market & why now                                                   | 10    |
| 09 | Where we are — done · doing · going                                    | 11    |
| 10 | Go-to-market — the five-phase plan                                     | 12–13 |

# What Bari is, in one breath

Bari turns the noise of a supermarket shelf into a single, defensible verdict. For any food or supplement category it answers the only question a shopper actually has: *“of the things in front of me, which is genuinely better — and why?”*

**499**

PRODUCTS SCORED

**13**

LIVE CATEGORIES

**10**

SCORING DIMENSIONS

**100%**

DIRECTLY VERIFIED DATA

## The gap we close

A shopper today has two bad options. Public databases and crowd apps are **unreliable** — wrong barcodes, missing fields, contributed errors. Brand marketing is, by design, **not neutral**. Israel's own government decided this matters: since 2020 it has mandated red warning labels on high-sugar, -salt and -fat products. But that signal is **binary** — a product either has a red label or it doesn't. Nobody tells an Israeli shopper, in Hebrew, *which* of two products is better and by how much. That continuous, explained, category-relative verdict is the whole of what Bari does.

### WHY IT'S CREDIBLE

- **Deterministic engine.** Scores come from auditable code, not an AI guessing — every score reproduces exactly from its own record.
- **Only verified data.** If a value wasn't read directly off the real product, the field is left empty and the page says so. We never fill gaps with guesses.
- **It isn't generous.** Across 499 products only 6% earn an A or better and 56% land at D or E. The engine grades honestly — it is not a marketing tool.

### WHY IT SCALES

- **A factory, not a craft.** Each category is the output of one repeatable pipeline — not a hand-built page.
- **One engine, every category.** Improve a rule once and all categories re-flow through the same machinery.
- **Hebrew-native.** The whole stack reasons over Hebrew ingredient text — a real, defensible moat in this market.

### POSITIONING

“Bari is the **category-relative, evidence-graded shelf** for Israeli food and supplements — the only place a shopper gets an honest verdict, with the reason shown, in their own language.”

# What's live, and what it tells a shopper

Bari is a working Hebrew web platform built around **comparison pages**. Each one takes a real shelf, scores *every product individually*, and ranks them so a shopper understands the shelf in 15–20 seconds on a phone.

13 CATEGORIES SHIPPED · 499 PRODUCTS

| CATEGORY                         | N   | STATUS |
|----------------------------------|-----|--------|
| Bread                            | 29  | LIVE   |
| Cheeses (yellow / hard / brined) | 112 | LIVE   |
| Milk & alternatives              | 18  | LIVE   |
| Hummus & spreads                 | 57  | LIVE   |
| Juices                           | 17  | LIVE   |
| Cereals · Granola                | 45  | LIVE   |
| Cookies · Cakes · Salty snacks   | 205 | LIVE   |
| Protein & snack bars             | 16  | LIVE   |
| Supplements (magnesium pilot)    | —   | BUILD  |

Grouped rows are broken into individual categories in Figure 2 (p.7). All live categories pass a structural conformance check — a change to the scoring engine re-flows cleanly through every one of them.

THE ENGINE GRADES HONESTLY — FULL CORPUS



## What a comparison page gives the shopper

THE VERDICT

A score (0–100) and grade (S–E) for *each* product, ranked within its category — no vague verbal labels layered on top of the number.

THE REASON

A two-line human verdict per product — where it stands, why, and the catch. The real driver is named (calorie density, sodium, fat source...).

THE HONESTY

A category caveat, a confidence state when data is thin, and an explicit “data could not be retrieved” instead of a guessed value.

## How a product becomes a score

Every product moves through three stages — **extract** → **understand** → **judge**. Real retail data in; an individual, defensible verdict out. The judging stage never sees an outside database; it only sees what was directly captured.




Figure 1 — The pipeline. The engine never reads an external database; it only sees what the scrape captured.

### EXTRACT

A four-step scraper pulls product pages from up to four Israeli chains; an OCR step reads physical label photos. The output is raw, per-product observations. If a field isn't found, it stays empty — there is no imputation anywhere in the system.

### UNDERSTAND

Software reads the Hebrew ingredient list — white vs. whole-grain flour, fermentation, additives and how many, the protein source, the food structure — and writes one clean, canonical record per product, tagged with how confident it is.

### JUDGE

The scoring engine turns that record into a score and grade and writes a full reasoning record alongside it. The website only ever displays this output, so the number on screen always traces back to the record that produced it.

### THE IMPORTANT SEPARATION

AI helps *read* messy Hebrew labels and *write* the consumer copy. It never sets a nutrition value or a score — that is done by deterministic code. This is exactly what makes Bari's output reproducible and defensible, where an “AI nutrition app” is not.

# How a score is computed

The engine reads a product across **ten dimensions**, then runs four discipline steps whose entire job is to make sure no single problem is counted twice. It is plain, deterministic code — same product, same score, every time.

## THE TEN DIMENSIONS (EACH 0–100)

| DIMENSION                       | WEIGHT |
|---------------------------------|--------|
| Processing quality              | 15%    |
| Nutrient density                | 15%    |
| Calorie density                 | 15%    |
| Glycemic quality                | 12%    |
| Protein quality                 | 10%    |
| Additive quality                | 10%    |
| Fat quality                     | 8%     |
| Satiety support                 | 6%     |
| Regulatory quality (red labels) | 5%     |
| Whole-food integrity            | 4%     |

## THEN FOUR DISCIPLINE STEPS

**Guardrails**  
hard ceilings for serious problems (e.g. trans fat)

**Hyper-palatability check**  
flags engineered fat-sugar / fat-salt combinations

**De-duplication step**  
stops one root problem penalising the score twice

**Confidence ceiling**  
caps the score when the data is thin

**final score 0–100 · grade S–E**

grades reflect the real spread, not an arbitrary curve

### “DON'T TEN DIMENSIONS DOUBLE-COUNT?” — THE MOST IMPORTANT QUESTION, ANSWERED

It's the right thing to ask: a red label is *triggered* by sugar, salt or fat — which also live in the glycemic, regulatory and fat dimensions. Two deliberate design choices prevent double-counting. **First, weighting:** the overlapping dimensions are intentionally light — regulatory quality is only 5%, whole-food integrity 4% — so they nudge, they don't dominate. **Second, a de-duplication step** exists for exactly this: when several rules trace back to one root cause, the engine keeps the primary signal at full weight and demotes the echoes, with a cap per root cause. **Worked example:** a sugary cereal loses points on glycemic quality — its primary flaw — and would *also* trip the red-label dimension. The de-duplication step recognises both trace to sugar, keeps the glycemic penalty in full, and shrinks the regulatory echo to a fraction. Sugar is paid for once, not twice. (We're happy to walk a technical partner through the implementation under NDA.)

### EVERY PRODUCT GETS ITS OWN SCORE — AND THEY GENUINELY DIFFER

Scoring is per-product, not per-category. The chart on the next page shows it: within a single category, products span a wide range — yellow cheeses run from 23 to 87, milks from 34 to 85. The engine separates the genuinely better product from the worse one on the same shelf; it does not hand a category one verdict.

## Why the numbers can be trusted

A score is only as good as the discipline behind it. Three rules exist so that any verdict can be defended publicly — to a shopper, a brand, a regulator, or a journalist. And the proof is in the spread: the engine differentiates, and it isn't generous.

### 1 · ONLY VERIFIED DATA

The only source is the product itself, read directly. Crowd-sourced databases are never used — not for nutrition, ingredients, names, or images. If we didn't capture it, the page says “data could not be retrieved.” Unknown is acceptable; a guess is not.

### 2 · EVERY CLAIM IS SOURCED

Every scoring rule and every consumer claim is tied to a recorded source. Rules can't quietly accumulate — each must be justified, scoped, and reversible. Vague provenance is rejected before anything ships.

### 3 · NOTHING INVENTED

The system never fabricates a product, a nutrition value, or an ingredient. A product with missing data is dropped, not guessed. Honesty over coverage — a smaller true shelf beats a larger fake one.

## The proof: real spread within every category




Figure 2 — Min-mean-max score per category (real corpus). Wide bars = the engine separates products; low means in cookies/cakes/snacks = it refuses to flatter junk.

## How it's built

Bari is two clean halves that meet at one place. A **data & scoring workshop** does all the thinking and produces finished, frozen results; a **website** simply displays them. Because the two only meet through finished files, the published number can never drift from the reasoning that produced it.

### THE WORKSHOP — WHERE THE THINKING HAPPENS

A collection of automated tools that gather and judge product data. In order:

- **Collectors** visit retailer websites and read product pages; a second tool reads photographs of physical labels.
- **A Hebrew-label reader** interprets the ingredient list — flour type, additives, fermentation, protein source.
- **The scoring engine** applies the ten dimensions and discipline steps and writes a finished result for each product.
- **Automatic tests** re-check the work so a change can't quietly break a past result.

### THE WEBSITE — WHERE SHOPPERS SEE IT

A fast, modern, mobile-first Hebrew site built right-to-left. It is deliberately “simple-minded”:

- It reads the finished results and **displays** them — nothing more.
- It never re-calculates, re-sorts or re-interprets a score; the ranking is decided upstream.
- Every page follows one shared visual template, checked for contrast and Hebrew correctness.
- Adding or updating a category means re-running the workshop, not editing pages by hand.

### WHY THIS DESIGN IS THE ASSET

Because all the intelligence is finished and frozen *before* the website sees it, three things follow: the site can never show a number that disagrees with its own reasoning; the engine and the website can be improved independently; and a whole category re-generates with a single command, not by re-doing work. That is what makes Bari a **platform**, not a pile of pages — and what makes new categories cheap to add.

Under the hood, for completeness: the workshop is Python; the website is a modern React/Next.js application in TypeScript. The two communicate through generated data files — the single, deliberate seam between them.



## The team that produces it

Bari runs as a small **studio of specialists** — a set of focused roles, each owning one job, coordinated by a single lead and held to account by an independent challenger. Data collection, research, content and frontend work are AI-assisted, which is how a lean operation produces a national-scale catalogue — *but the scoring logic itself is deterministic code, not an AI judgment.*

### DATA

Runs the collection and scoring pipeline — scrapes the shelves, manages the corpus, produces the scored results for each category.

### NUTRITION

Owns the scoring philosophy and food-science reasoning — what each dimension means, how categories are interpreted, why a product scores as it does.

### RESEARCH

Gathers the evidence behind every rule and claim — scientific literature, supplement evidence, competitor and market analysis. Produces facts; doesn't decide.

### PRODUCT

Makes the strategic calls — what to build, what to cut, which category is next, and the final go / no-go before a category goes live.

### CONTENT

Writes the consumer-facing words — the verdicts and explanations — in a consistent, honest, human Hebrew voice.

### FRONTEND & DESIGN

Builds and visually polices the website — implementing pages and enforcing one consistent, accessible, correct Hebrew layout.

### MARKETING

Owns growth once a category is live — search visibility, content, audience and distribution. (The engine of the go-to-market plan.)

### INDEPENDENT QA / CHALLENGER

The independent challenger. Tries to break every score and every claim before launch — can a verdict be defended in public? It raises problems; it cannot approve its own fixes.

### LEAD / COORDINATOR

The founder's operational role: coordinates all of the above off a single shared task list — assigns the work, checks each result against the actual evidence, and is the only role that can mark a job “done.” No “trust me, it's finished.”

### THE RULE THAT PROTECTS TRUST

Nothing a shopper will read ships without **two** independent sign-offs — the Content role *and* the independent challenger. Every consumer sentence, every category, every time.

## The market & why now

Three forces line up behind Bari: a government that has already made nutrition transparency the law, a fast-growing health-and-supplements market, and a shift in how people find trustworthy answers online. Bari is built precisely for this moment.

~2,200

DEATHS/YR THE RED-LABEL REFORM  
AIMS TO PREVENT

~600M

ANNUAL HEALTHCARE COST IT AIMS  
TO SAVE

~\$0.5B

ISRAEL VITAMINS & MINERALS  
MARKET BY 2029

### 1 · THE STATE ALREADY VALIDATED THE PROBLEM

Since 2020 Israel mandates red warning labels on products high in sugar, salt or saturated fat — a policy the Health Ministry projects will cut consumption ~30%, prevent ~2,200 deaths a year and save ~600M annually. The state has decided point-of-purchase nutrition guidance matters. But its signal is **binary**. **Bari is the continuous, comparative layer on top of it** — not “does this have a red label,” but “which of these is actually better, and by how much.”

### 2 · A GROWING, HEALTH-CONSCIOUS MARKET

Supplements are the second-largest over-the-counter category in Israel after painkillers, and the vitamins & minerals market alone is projected to approach **US\$0.5 billion by 2029**, growing ~5% a year. Consumers are actively shifting to health-conscious buying — and increasingly shopping online, where a comparison layer is most useful.

### 3 · THE SEARCH SHIFT — WHY CREDIBILITY WINS NOW

In 2026, AI-generated answers sit above search results and absorb the clicks. The sites that still win are the ones AI *cites* — and citation now flows to **original data** and **credentialed authorship** (the E-E-A-T signals search engines now reward). Bari is exactly that profile: a proprietary 499-product dataset, plain honest verdicts, and — with the right partner — a licensed clinical authority behind the methodology. The integrity discipline isn't overhead; in this search era it is the growth strategy.

Sources: Israel Ministry of Health / EfshariBari front-of-package reform — the ~30% / ~2,200-deaths / ~600M figures are the MoH's *pre-reform projections* (2019–20), not measured outcomes. · Euromonitor / Statista — Dietary Supplements & Vitamins-and-Minerals, Israel. · Search Engine Land & 2026 SEO benchmark studies (AI-Overview citation & E-E-A-T authorship). Figures are directional market indicators, not a Bari forecast.

# Done · doing · going

The hard, defensible part — a trustworthy engine and a repeatable factory — is built. What remains is breadth, audience and revenue: the phase where the right partner adds the most.




Figure 3 — The engine and factory are done; focus shifts to breadth, audience and revenue.

## The plan: five phases

A deliberate sequence. Each phase earns the right to the next: trust builds audience, audience builds an irreplaceable dataset, the dataset opens institutional doors, institutions justify funding, and funding opens new markets and models. Revenue stays modest and clean early — credibility is the asset we are compounding.

1

### Audience, traction & credibility

Win attention and trust. Comparison pages are durable assets that answer “which X is healthier” for years. The 2026 playbook: rank by being **cite-worthy** — original data and a credentialed authority — so AI answers reference Bari; own the Hebrew long-tail of category questions; build a newsletter as a direct, owned audience; grow on YouTube/Reddit-style channels where AI citations concentrate. Monetisation here is light and non-intrusive: modest referral and sponsorship, never over the top, never touching a score. **In parallel, in the pipeline:** begin a personalised consumer engine that turns the scored catalogue into “best for *you*” given goals and restrictions — the feature that converts visitors into a returning, logged-in audience.

2

### Become *the* trusted database — and sell it where it's needed

The audience phase produces the real moat: a clean, verified, continuously-scored national food & supplement dataset. That asset is valuable to the institutions that lack it. **Health funds (Kupot Holim):** all four run advanced member apps — several already with nutrition and healthy-lifestyle features — and serve the whole population. Bari can power an embedded “what should I actually buy” layer inside those apps (e.g. a Clalit or Maccabi member scanning a shelf and getting a trusted verdict). **Government & public health:** a natural complement to the red-label reform — the continuous, comparative data the binary label can't provide. **Retail chains:** private-label benchmarking and a healthier-shelf signal. This is where meaningful, defensible B2B revenue begins.

#### WHY PHASES 1 → 2 ARE ONE MOTION

You can't sell a “trusted national database” to a health fund until the public already treats you as trustworthy. Phase 1 manufactures exactly the credibility and the dataset that make Phase 2 a credible conversation. They are not two businesses — they are one compounding loop.

## Phases 3–5, and the flywheel

3

### Clinical legitimacy — the licensed-professional path

Dietitians in Israel are a regulated, licensed profession (Ministry of Health registration and exam). Evaluating Bari's fit with a licensed clinical authority — for methodology sign-off and professional-facing tools — converts trust into formal credibility. It strengthens the institutional sale in Phase 2, satisfies the “credentialed author” signal that wins modern search, and opens a professional channel (clinics, dietitians citing and using Bari). **This is the most direct point of overlap with a nutritionist partner.**

4

### Funding & international expansion

With audience, an institutional dataset business, and clinical legitimacy, Bari is fundable. The same engine is country-agnostic — the moat is the *local* label data and language, exactly the part Bari already proved it can build. Capital funds replication into new markets, starting where regulatory tailwinds (front-of-pack labelling) and health-conscious demand mirror Israel's.

5

### Diversify into commerce

Once Bari is where people decide *what* to buy, helping them buy it is the natural next step — contextual, opt-in commerce: “add the better option to your cart,” healthier-basket building, and curated marketplace links, all kept strictly downstream of and walled off from the score. A monetisation of intent, not of influence.

## The flywheel that ties it together




Figure 4 — Trust earns audience; audience builds the dataset and revenue; revenue funds more categories, which deepen trust. The factory makes the loop cheap to spin.

Channel research: Israeli HMO digital-app landscape (Clalit, Maccabi, Meuhedet, Leumit); MoH dietitian licensing; 2026 SEO / newsletter / contextual-commerce practice. See Section 08 sources.

IN ONE SENTENCE

**The engine and the factory are built, honest, and live. The opportunity now is trust → audience → an irreplaceable national dataset → institutions, funding and new markets.**

---

**Tom Bar Haim** — Founder & Owner, Bari

tbarhaim@gmail.com

Prepared 21 June 2026 · Confidential